

Process synchronization

Revy Naik V
Assistant Professor
Computer science & engineering
RGUKT Basar

objectives

- To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data.
- To present both software and hardware solutions of the critical-section problem.
- To introduce the concept of an atomic transaction and describe mechanisms to ensure atomicity.

Cooperating processes concurrent access

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. Initially, count is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

producer-consumer

Producer:

```
while (true) {  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

Producer-consumer

Consumer:

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
}
```

Race-condition

- **count++** could be implemented as

```
register1 = count
register1 = register1 + 1
count = register1
```
- **count--** could be implemented as

```
register2 = count
register2 = register2 - 1
count = register2
```
- Consider this execution interleaving with “count = 5” initially:
 - S0: producer execute **register1 = count** {register1 = 5}
 - S1: producer execute **register1 = register1 + 1** {register1 = 6}
 - S2: consumer execute **register2 = count** {register2 = 5}
 - S3: consumer execute **register2 = register2 - 1** {register2 = 4}
 - S4: producer execute **count = register1** {count = 6}
 - S5: consumer execute **count = register2** {count = 4}

The Critical section problem

- Consider a system consisting of n processes $\{P_0, P_1, \dots, P_{n-1}\}$. Each process has a segment of code, called a critical section, in which the process may be changing common variables, updating a table, writing a file, and so on.
- When one process is executing in its critical section, no other process is to be allowed to execute in its critical section.

```
do{  
    entry section;  
    Critical section;  
    exit section;  
    remainder section  
}while(true);
```

Critical section

- The critical section problem is to design a protocol that the processes can use to cooperate.
- Each process must request permission to enter its critical section. The section of code implementing this request is the **entry section**.
- The critical section may be followed by an exit section.
- The remaining code is the remainder section.

Solution to Critical-Section Problem

Requirements:

1. Mutual Exclusion - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.
 2. Progress - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely.
 3. Bounded Waiting - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.
- Assume that each process executes at a nonzero speed
 - No assumption concerning relative speed of the N processes

Simple lock variable

Initial lock value is zero

```
lock=0;
```

```
do{
```

```
while(lock!=0);
```

```
    Lock=1; entry section
```

```
    Critical section;
```

```
    Lock = 0 ; exit section
```

```
}while(1);
```

Strict alternation

Process P_i :

Non critical section ;

while(turn!=i); //entry section process p_i enters into C.S

Critical section;

turn=j; /*Exit section //process p_i is holding process p_j critical section

Process P_j

Non- critical section;

while(turn!=j); //entry section process p_j enters into C.S

Critical section;

turn=i; //Exit section process p_j holding process p_i C.S

Progress is not guaranteed because one process holding another process critical section.

Hardware solution

Test And Set Instruction

Definition:

```
boolean TestAndSet (boolean  
*target)  
{  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

Must be executed atomically

Hardware solution

Shared Boolean variable lock, initialized to false.

Solution:

```
do {  
    while ( TestAndSet (&lock ))  
        ; // do nothing  
        // critical section  
    lock = FALSE;  
        // remainder section  
} while (TRUE);
```

Peterson's solution

- Two process solution
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:

int **turn**;

Boolean **flag**[2]

- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag**[i] = true implies that process P_i is ready!

Process P_i

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    critical section  
    flag[i] = FALSE;  
    remainder section  
} while (TRUE);
```

Process Pj

```
do {  
    flag[j] = TRUE;  
    turn = i;  
    while (flag[i] && turn == i);  
    critical section  
    flag[j] = FALSE;  
    remainder section  
} while (TRUE);
```


Semaphore

- Synchronization tool that does not require busy waiting
- Semaphore S – integer variable
- Two standard operations modify S : `wait()` and `signal()`
Originally called `P()` and `V()`
Also called `down()` and `up()`
- Less complicated
- Can only be accessed via two indivisible (atomic) operations. Two types of semaphores: binary semaphore (range is 0 to 1) and counting semaphore (range is -infinity to +infinity).
- (-ve) value indicates no of processes blocked.

Counting semaphore

- We know the no of processes that are blocked or suspended.

```
Down (semaphore s){  
    s.value=s.value-1;  
    if(s.value<0){  
        Put process PCB in suspend list sleep();  
    }  
    Else{  
        Return;  
    }  
}
```

Semaphore (up operation)

```
Up (semaphore s){  
    s.value=s.value+1;  
    if(s.value<=0){  
        Select a process from suspend list;  
        Wakeup();  
    }  
}
```

Counting semaphore

- In a certain computation the value of a counting semaphore (S) is initialized to 12, the following up & down operations are performed in the given order. 15P, 12V, 8P, 6V, 2P, 3V, 8P, 4V what is the current value of semaphore?

Binary semaphore

```
Down (semaphore s){  
    if(s.value==1){  
        s.value=0;  
    }  
    Else{  
        Block this process and place in suspend list;  
        Sleep();  
    }  
}
```

Binary semaphore

- We don't know the no of blocked processes as long as there is any blocked process the value of S remains at 0.

```
up(semaphore s){  
    If(suspend lis is empty)  
    {  
        s.value=1;  
    }  
    Else {  
        Select process from suspend list;  
        Wakeup();  
    }  
}
```

Binary semaphore

- Let B semaphore $s=1$, 10p, 8v, 16p, 9v, 2p, 4v operations then what is the value of S and length of queue S.L()?

9	1	17	8	10	6
0	0	0	0	0	0

value of s is 0 and queue length is 6.

Bounded buffer problem using semaphore

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N .
- Semaphore `empty=N` // no of empty slots is the buffer
- Semaphore `full =0` // no of full slots in the buffer
- Semaphore `mutex=1` // used to ensure producer mutual exclusion between producer P & consumer C processes on the buffer.

Producer process

Void producer (void){

Int item p

While(1){

produce_item(item p):

down(empty) // when buffer is full empty = 0 and producer is blocked

down(mutex)

buffer[in]=item p;

In =(in+1) % n;

up(mutex);

up(full);

}

}

Consumer process

```
Void consumer (void){  
    Int item c;  
    While(1){  
        down(full) // when buffer is free full = 0 so consumer will be blocked  
        down(mutex)  
        Int c=buffer[out]  
        out=(out+1) % n  
        up(mutex);  
        up(empty);  
        process_item(item c)  
    }  
}
```

Readers and writers problem

- 1R-1W (one reader one writer)
- NR-1W (multiple readers – one writer)
- NR-MW (multiple readers -writers)
- `Int rc = 0 // readers count in dbms`
- Semaphore `mutex = 1 // used between readers to access 'rc' in a mutual exclusion manner.`
- Semaphore `db=1 // used between readers & writers to access DBMS in a M.E manner.`

Readers - writers

```
Void reader (void){  
    While(1){  
        down(mutex);  
        rc=rc+1;  
        if(rc==1) down (db);  
        up(mutex)  
        Read DBMS  
        down(mutex)  
        rc=rc-1;  
        if(rc==0) up(db);  
        up(mutex)  
    }  
}
```

readers-writer

```
Void writer(void){
```

```
    While(1)
```

```
    {
```

```
        down(db);
```


```
        <write DBMS>
```

```
        up(db);
```

```
    }
```

```
}
```

Dining philosopher



Dining philosopher

Do{

Wait(chopstick[i]);

wait(chopstick[(i+1)%5]);

....

//eat

....

signal(chopstick[i]);

signal(chopstick[(i+1)%5]);

....

//think

....

}while(true);

Monitors


- A high level abstraction that provides a convenient and effective mechanism for process synchronization.
- A monitor type presents a set of programmer defined operations that provides mutual exclusion with the monitor.
- A procedure defined within a monitor can access only those variables declared locally within the monitor and its formal parameters.
- Similarly the local variables of a monitor can be accessed by only the local procedures.

Monitor

- The monitor construct ensures that only one process at a time can be active within the monitor.
- Condition construct condition x, y ; the only operations that can be invoked on a condition variable are `wait()` and `signal()`;
- The operation `x.wait()`; mean that the process invoking this operation is suspended until another process invokes `x.signal()`;
- The `x.signal()` operation resumes exactly one suspended process.

Monitor

Schematic View of a Monitor



monitor

```
Monitor  monitor name
{
    //shared variable declarations
    Procedure P1(.....){
        .....
    }
    Procedure P2(.....){
        .....
    }
    .
    .
    .
    Procedure Pn(.....){
        .....
    }
}
```

Dining philosophers solution using monitor

`enum{THINKING, HUNGRY, EATING} state[5];`

- Philosopher i can set the variable `state[i]=EATING` only if her two neighbors not eating: `(state[(i+4)%5]!=EATING)` and `(state[(i+1)%5]!=EATING)`.
- Condition `self[5]`; in which philosopher i can delay herself when she is hungry but is unable to obtain the chopsticks she needs.

DP solution using monitor

```
Monitor dp{  
    enum{THINKING, HUNGRY, EATING} state[5];  
    Condition self[5];  
    Void pickup(int i){  
        state[i]=HUNGRY;  
        test(i);  
        If(state[i]!=EATING)  
            self[i].wait();  
    }  
}
```

DP solution using monitor

```
Void putdown(int i){
    state[i]=THINKING;
    test((i+4)%5);
    test((i+1)%5);
}

Void test(int i){
    if((state[(i+4)%5]!=EATING) && (state[i]==HUNGRY) && (state[(i+1)%5]!=EATING))
    {
        state[i]=EATING;
        Self[i].signal();
    }
}
```

DP solution using monitor

```
initialization_code(){  
    for(int i=0; i<5; i++){  
        state[i]=THINKING;  
    }  
}  
}
```